

Manjunath Popuri

+1 (917) 815-9580 · manjunathpopuri2@gmail.com · [LinkedIn](#) · Binghamton, NY (Open to Relocation)

SUMMARY

Data Scientist with 3+ years of combined research, industry, and academic experience applying statistical methods, machine learning, and cloud-native AI solutions to deliver measurable business outcomes. Skilled in A/B testing, Bayesian inference, predictive modeling, NLP (BERT, Transformers), and data visualization. Proven ability to translate business problems into data-driven experiments, build scalable ML pipelines on Azure & AWS, and communicate insights through executive storytelling and dashboards.

EDUCATION

University at Binghamton, State University of New York, Binghamton, NY, USA <i>Masters in Computer Science (Specialization: AI)</i>	AUG 2023 - MAY 2025
Vasireddy Venkatadri Institute of Technology, Guntur, Andhra Pradesh, India <i>Bachelors in Computer Science and Engineering</i>	AUG 2019 - JUN 2023

TECHNICAL SKILLS

Languages: Python, R, Java, SQL, C/C++, JavaScript

ML & AI: TensorFlow, PyTorch, Scikit-learn, OpenCV, Apache Spark MLlib, A/B testing, Hugging Face, Transformers, Generative AI, LLM Fine-tuning, NLP (BERT, TF-IDF, Word2Vec)

Data Tools: Pandas, NumPy, Tableau, Power BI, Matplotlib, Seaborn, SpaCy, PySpark

Cloud & MLOps: AWS, Azure, Docker, CI/CD Pipelines, Git

Databases: MySQL, PostgreSQL, MongoDB, MS SQL

WORK EXPERIENCE

Binghamton University, Remote, Binghamton, NY, USA	Research Intern	JUL 2025 - PRESENT
<ul style="list-style-type: none">Developed a quantized LLM (Mistral-7B, Phi-3) for predictive expense classification (90%+ accuracy), optimized for low-resource devices.Built ML pipelines in Python, PyTorch, Hugging Face with feature engineering, time-series modeling, and data wrangling of financial transactions.Designed and deployed prescriptive ML models on AWS SageMaker, providing personalized budget recommendations.Delivered interactive dashboards (Matplotlib, Seaborn, Power BI) with business storytelling, enabling non-technical stakeholders to make data-driven decisions.		
Binghamton University, Hybrid, Binghamton, NY, USA	Graduate Research & Teaching Assistant	AUG 2023 - MAY 2025
<ul style="list-style-type: none">Applied A/B testing, Bayesian inference, and quasi-experimental methods to validate NLP models analyzing Russell 3000 earnings calls, boosting predictive accuracy by 18%.Designed ML pipelines with BERT, Transformers, Scikit-learn, and deployed at scale using Azure ML.Conducted hypothesis testing and regression analysis on financial data to ensure statistical rigor in insights shared with faculty/industry.Created executive-level dashboards (Power BI, Tableau) with UX design principles and mentored 150+ students in ML and statistical experimentation.		
Cognizant Technology Solutions, Onsite, Hyderabad, India	Data Scientist	JUL 2022 - JUN 2023
<ul style="list-style-type: none">Delivered predictive modeling and classification solutions for healthcare/finance clients, improving accuracy by 30% with Scikit-learn, TensorFlow, and PyTorch.Applied A/B testing, causal inference, and Bayesian optimization to fine-tune ML models, reducing false positives by 22%.Built scalable ETL pipelines with PySpark and SQL for big-data processing on millions of records.Deployed production-ready models via Azure ML & AWS SageMaker with Docker and CI/CD pipelines (GitHub Actions) and presented insights through Tableau dashboards.		

PROJECTS

Advanced Movie Recommendation System (Unsupervised Learning): Built a content-based recommendation engine using TF-IDF, Word2Vec, PCA, and K-Means, boosting semantic accuracy by 20% and scaling similarity computations with parallel processing.

Tech Stack: Python, Scikit-learn, NLTK, TF-IDF, Word2Vec, PCA, K-Means clustering

Real-Time Digit Recognition App (Deep Learning): Developed Android app for real-time handwritten digit recognition using MobileNetV3 and ResNet-50, achieving <100ms latency and deployed via Firebase ML Kit, enabling fast, GPU-accelerated predictions.

Tech Stack: Python, TensorFlow, Keras, TensorFlow Lite, Android Studio, Java, Firebase ML Kit, OpenCV, NumPy

CERTIFICATION

- Generative AI for LLM's, AWS Cloud Foundations & Architecting**